

**Question 1.** [12 MARKS]

Consider this schema for Twitter, a social media platform where users post messages called “tweets”.

**Relations**

User(userID, name, email)

A Twitter user.

Tweet(tweetID, userID, content, day)

The user with *userID* made a tweet containing *content* on *day*.

Follows(a, b)

User *a* follows user *b* on Twitter, which means that

*a* has subscribed to *b*’s tweets.

Likes(who, what, d)

User *who* liked tweet *what* on day *d*.

**Integrity Constraints**

$\text{Tweet}[\text{userID}] \subseteq \text{User}[\text{userID}]$

$\text{Follows}[a] \subseteq \text{User}[\text{userID}]$

$\text{Follows}[b] \subseteq \text{User}[\text{userID}]$

$\text{Likes}[\text{who}] \subseteq \text{User}[\text{userID}]$

$\text{Likes}[\text{what}] \subseteq \text{Tweet}[\text{tweetID}]$

**Part (a)** [2 MARKS]

Does the schema enforce this constraint: A user cannot like the same tweet twice? Circle one:

**Solution:**

☐ Yes

☐ No.

Since who and what together form a key for relation Likes, the same user-tweetID combination cannot occur more than once in relation Likes.

**Part (b)** [2 MARKS]

Does the schema enforce this constraint: You can’t follow yourself? Circle one:

**Solution:**

☐ Yes

☐ No.

This would enforce it:

$$\sigma_{a=b} \text{Follows} = \emptyset$$

**Part (c)** [1 MARK]

Suppose relation *Likes* has 300 tuples. How many tuples could *Users* have? Circle all that apply:

**Solution:**

☐ 0

☐ 1

☐ 256

☐ 300

☐ 912

**Part (d)** [2 MARKS]

Suppose relation *User* has  $m$  tuples and relation *Tweet* has  $n$  tuples. What is the maximum number of tuples that relation *Likes* can have?

**Solution:**

$$m \times n$$

Explain how the schema imposes this limit:

**Solution:**

Each user-TweetID combination can only occur once, and the total possible number of combinations is  $m \times n$ .

**Part (e)** [3 MARKS]

Suppose we add the following constraint:  $\text{Likes}[\text{who}] \subseteq \text{Follows}[\text{b}]$ . Make the smallest possible non-empty instance of relations *Likes* and *Follows* that violates this constraint:

**Solution:**

Likes:	who	what	d
	miriam	T23	Jan 1, 2016

Follows:	a	b
	drizzy	dianeh

Express this constraint in English:

**Solution:** You can't like any tweets unless you have follower(s).

What kind of constraint is it? Circle all that apply:

**Solution:**
☐ referential integrity constraint

☐ foreign-key constraint

☐ integrity constraint

**Part (f)** [2 MARKS]

Which of the following queries can be expressed using the same form of relational algebra that we used in class and on Assignment 1, that is: the operators  $\Pi, \sigma, \bowtie, \bowtie_{condition}, \times, \cap, \cup, -, \rho$  and assignment? Circle all that apply.

**Solution:**

1. **Yes:** Find everyone who follows 6 or more people who have never liked a tweet.
2. **No:** Let's say user X is "upstream" of Y if either X follows Y, or X follows someone else who is upstream of Y. Find every user who is upstream of the person with userID 'Oprah'.
3. **Yes:** Find the second last tweet from the person with userID 'Oprah'.
4. **No:** Find the user who follows the most people.
5. **Yes:** Find the user who made the first tweet.

**Question 2.** [8 MARKS]

Here is the schema from Assignment 1. A few attributes and relations have been omitted for simplicity.

**Relations**

Product(DIN, manufacturer, name, form, schedule)

A tuple in this relation represents a drug product.

Price(DIN, price)

The price of a drug product.

Prescription(RxID, date, patient, drug, doctor)

A prescription for *drug* was written on *date* for *patient* by *doctor*. Attribute *patient* is the patient's OHIP number.

Filled(RxID, date, pharmacist)

Prescription *RxID* was filled by *pharmacist* on *date*.

Attribute *pharmacist* is the pharmacist's OCP number.

**Integrity constraints**

Price[DIN]  $\subseteq$  Product[DIN]

Prescription[drug]  $\subseteq$  Product[DIN]

Filled[RxID]  $\subseteq$  Prescription[RxID]

$\Pi_{\text{schedule}} \text{Product} \subseteq \{\text{"prescription"}, \text{"narcotic"}, \text{"OTC"}\}$

Write a query in relational algebra to report the OHIP number of every patient who has had a prescription that (a) was for the most expensive drug product (or a product tied for most expensive) and (b) they never filled.

Use only the basic operators  $\Pi, \sigma, \bowtie, \times, \cap, \cup, -, \rho$ , and assignment.

**Solution:**

– DIN of a drug that is not the most expensive drug.

$NotMax(DIN) := \Pi_{P1.price < P2.price}(\rho_{P1}Price \times \rho_{P2}Price)$

– DIN of a drug that IS the most expensive drug.

$Max(DIN) := (\Pi_{DIN}Price) - (\Pi_{DIN}NotMax)$

– This prescription for this patient is for the most expensive drug, or for a drug that is tied for most expensive.

$MaxPrescription(RxID, patient) := \Pi_{RxID, patient}(\sigma_{drug=MAX}Prescription \times Max)$

– This patient has had a prescription for the/a most expensive drug that they never filled.

$Answer(patient) := \Pi_{patient}[(\Pi_{RxID}MaxPrescription) - (\Pi_{RxID}Filled)]$

*Continue your answer here if more space is needed.*

**Question 3.** [6 MARKS]

Suppose we have implemented the Twitter schema from Question 1 in SQL, and the tables currently contain the following:

Profile:

userid	name	email
adele	Adele Adkins	
drizzy	Drake	
potus	Barack Obama	potus@gov.us
rjm	Renee Miller	rjm@cs

Follows:

a	b
potus	drizzy
drizzy	rjm

Tweet:

tweetid	userid	content	day
15	adele	Hello	2016-10-16
61	adele	It's me	2016-10-16
33	potus	6 weeks	2016-10-11
28	rjm	in the 6	2016-10-10

Likes:

who	what	d
drizzy	61	2016-10-18
rjm	33	2016-10-17
drizzy	15	2016-10-16
potus	15	2016-10-16

Show the output of each of the queries below. If a query will not run successfully, write “Illegal”.

**Solutions**

----- (1)

```
SELECT who
FROM Likes JOIN Tweet ON what = tweetID
WHERE userID = 'adele';
```

```
who
-----
drizzy
drizzy
potus
(3 rows)
```

----- (2)

```
SELECT userID, count(tweetID), count(day)
FROM Tweet
GROUP BY userID;
```

```
userid | count | count
-----+-----+-----
rjm    | 1     | 1
adele  | 2     | 2
potus  | 1     | 1
(3 rows)
```

----- (3)

```
SELECT count(*) AS num1, count(email) AS num2
FROM Profile;
```

num1	num2
4	2

(1 row)

----- (4)

```
SELECT name, content
FROM Profile NATURAL RIGHT JOIN Tweet;
```

name	content
Adele Adkins	Hello
Adele Adkins	It's me
Barack Obama	6 weeks
Renee Miller	in the 6

(4 rows)

----- (5)

```
(SELECT a AS userID
FROM Follows
WHERE b = 'drizzy')
UNION ALL
(SELECT userID
FROM Tweet
```

userid
potus
potus
rjm

(3 rows)

----- (6)

```
SELECT tweetID, count(who)
FROM Tweet, Likes
WHERE tweetID = what;
```

psql:questions.sql:33: ERROR: column "tweet.tweetid" must appear in the GROUP BY clause or be used in an aggregate function  
LINE 1: SELECT tweetID, count(who)

**Question 4.** [4 MARKS]

Write a query to find the userID of everyone who has made more than one Tweet. Ensure that it would work on any instance of the database, not simply the one above.

**Solution:**

This approach would work just as well if we were restricting to people who have made more than 100 Tweets:

```
SELECT Profile.userid
FROM Profile JOIN Tweet ON Profile.userID = Tweet.userID
GROUP BY Profile.userID
HAVING count(*) > 1;
```

This approach would work for the question, but would not scale well to larger cutoffs!

```
SELECT DISTINCT Profile.userid
FROM Profile, Tweet t1, Tweet t2
WHERE t1.userID = t2.userID and t1.tweetID <> t2.tweetID and Profile.userID = T1.userID;
```



